

TENSORFLOW EXAMPLE: RECURRENT NEURAL NETWORK

Problem: Adding two 8-bit numbers together.

For example

$$\begin{array}{r} 00001001 + 00111100 = 01000101 \\ 9 \quad + \quad 60 \quad = 69 \end{array}$$

```
In [1]: import numpy as np
import tensorflow as tf
```

```
In [2]: # Define the dataflow graph
time_steps = 8          # time steps which is the same as the length of the bit-string
input_dim = 2           # number of units in the input layer
hidden_dim = 16         # number of units in the hidden layer
output_dim = 1          # number of units in the output layer

# input X and target output Y
X = tf.placeholder(tf.float32, [None, time_steps, input_dim])
Y = tf.placeholder(tf.float32, [None, time_steps])

# define the RNN cell: can be simple cell, LSTM or GRU
cell = tf.nn.rnn_cell.BasicRNNCell(num_units=hidden_dim, activation=tf.nn.sigmoid)
# cell = tf.nn.rnn_cell.LSTMCell(hidden_dim, state_is_tuple=True)

# values is a tensor of shape [batch_size, time_steps, hidden_dim]
# last_state is a tensor of shape [batch_size, hidden_dim]
values, last_state = tf.nn.dynamic_rnn(cell, X, dtype=tf.float32)
values = tf.reshape(values, [time_steps, hidden_dim])

# put the values from the RNN through fully-connected layer
W = tf.Variable(tf.random_uniform([hidden_dim, output_dim], minval=-1.0, maxval=1.0), name='W')
b = tf.Variable(tf.zeros([1, output_dim]), name='b')
h = tf.nn.sigmoid(tf.matmul(values, W) + b)

# minimize loss, using ADAM as weight update rule
h_ = tf.reshape(h, [time_steps])
Y_ = tf.reshape(Y, [time_steps])
loss = tf.reduce_sum(-Y_ * tf.log(h_) - (1-Y_) * tf.log(1-h_))
train_step = tf.train.AdamOptimizer(0.1).minimize(loss)
```

```

In [3]: # Launch the graph
sess = tf.Session()
sess.run(tf.global_variables_initializer())

# Try out the example
# 00001001 + 00111100 = 01000101
# 9 + 60 = 69
x = np.array([
    # t=0 t=1 t=2 t=3 t=4 t=5 t=6 t=7
    [0,0], [0,0], [0,1], [0,1], [1,1], [0,1], [0,0], [1,0]
]).reshape([1,8,2])
y = np.array([0, 1, 0, 0, 0, 1, 0, 1]).reshape([1,8])

# train
sess.run(train_step, {X: x, Y: y})

# print result
[output_vals, _probs, _loss] = sess.run([values, h, loss], {X: x, Y: y})
print('Raw output values:')
print(output_vals)

# prediction
probs = np.array(_probs).reshape([8])
prediction = np.array([1 if p >= 0.5 else 0 for p in probs]).reshape([8])
print()
print('Probabilities: \n', probs)
print()
print('Prediction :', prediction)

# calculate absolute error
error = np.sum(np.absolute(y - probs))

# print out pre-training X-entropy Loss and absolute error
print('Absolute error : ', error)
print('X-entropy loss : ', _loss)

```

Raw output values:

```

[[ 0.47502092  0.47502097  0.52497905  0.47502092  0.52497905  0.52497888
   0.47502086  0.47502092  0.52497894  0.47502086  0.47502086  0.524979
   0.52497888  0.52497905  0.52497751  0.52497911]
 [ 0.14841537  0.61733449  0.75091577  0.28058618  0.56417257  0.68358868
   0.16095461  0.29389748  0.40222776  0.20322989  0.34588733  0.72239882
   0.79938602  0.78864682  0.58117068  0.52698869]
 [ 0.12746073  0.57632226  0.7341361  0.40259528  0.4287056  0.65959179
   0.15043941  0.31832516  0.43690103  0.18125151  0.38751271  0.717085
   0.80123001  0.73989534  0.50207788  0.4642967 ]
 [ 0.13945937  0.55163407  0.7223177  0.41603133  0.44143528  0.66948217
   0.15613405  0.33669403  0.43700132  0.18105952  0.3861483  0.70811087
   0.79458827  0.72580874  0.49339053  0.44845974]
 [ 0.13654141  0.48674026  0.80694008  0.30669942  0.56263697  0.65702188
   0.1532962  0.40506229  0.53345579  0.1348618  0.37143165  0.71475589
   0.78726411  0.70103151  0.52853388  0.38627633]
 [ 0.13166356  0.55130708  0.75171822  0.40902984  0.45188832  0.6567297
   0.15456353  0.31604448  0.44324529  0.17099875  0.38202119  0.69221586
   0.78805459  0.74973863  0.50217563  0.46639466]
 [ 0.17191656  0.56666446  0.76289117  0.33790958  0.57055444  0.66204184
   0.16484019  0.35837135  0.42414922  0.21562377  0.32464689  0.76766944
   0.80519629  0.74734473  0.57357621  0.56227982]
 [ 0.14956163  0.52249789  0.84678692  0.22361328  0.66431755  0.65289211
   0.15116327  0.40908679  0.51506007  0.15196115  0.30279168  0.78318781
   0.8053413  0.74328947  0.61467046  0.51279181]]

```

Probabilities:

```

[ 0.5623672  0.82815796  0.79375088  0.78649455  0.79683888  0.79787099
  0.82452643  0.84132642]

```

Prediction : [1 1 1 1 1 1 1 1]

Absolute error : 4.29662257433

X-entropy loss : 7.8703

In [4]: *# Let's train*

```
# generate the training data set
binary_dim = 8
largest_number = pow(2, binary_dim)
a_list = []
b_list = []
n_examples = 1000
for j in range(n_examples):
    a_list.append(np.random.randint(largest_number/2))
    b_list.append(np.random.randint(largest_number/2))

# train
n_epochs = 10
for i in range(n_epochs):

    for j in range(n_examples):
        a_int = a_list[j]
        b_int = b_list[j]
        c_int = a_int + b_int
        a = np.unpackbits(np.array([a_int], dtype=np.uint8))
        b = np.unpackbits(np.array([b_int], dtype=np.uint8))
        c = np.unpackbits(np.array([c_int], dtype=np.uint8))
        ab = np.c_[a,b]
        x = np.array(ab).reshape([1,binary_dim,2])
        y = np.array(c).reshape([1,binary_dim])
        sess.run(train_step, {X: x, Y: y})

    # print out loss for a random example
    if (i%1 == 0):
        # pick a random example out of the training data set
        k = np.random.randint(n_examples)
        a_int = a_list[k]
        b_int = b_list[k]
        c_int = a_int + b_int
        a = np.unpackbits(np.array([a_int], dtype=np.uint8))
        b = np.unpackbits(np.array([b_int], dtype=np.uint8))
        c = np.unpackbits(np.array([c_int], dtype=np.uint8))
        ab = np.c_[a,b]
        x = np.array(ab).reshape([1,binary_dim,2])
        y = np.array(c).reshape([1,binary_dim])

        # get predicted value
        [_probs, _loss] = sess.run([h, loss], {X: x, Y: y})
        probs = np.array(_probs).reshape([8])
        prediction = np.array([1 if p >= 0.5 else 0 for p in probs]).reshape([8])
        pred_int = np.sum(np.packbits(prediction))

        # calculate error
        error = np.sum(np.absolute(y - probs))

        print('Input 1      : ', a, ' (', a_int, ')')
        print('Input 2      : ', b, ' (', b_int, ')')
        print('True          : ', c, ' (', c_int, ')')
        print('Predicted       : ', prediction, ' (', pred_int, ')')
        print('Absolute error  : ', error)
        print('X-entropy loss  : ', _loss)
        print('-----')
        print()
```

```
Input 1      : [0 0 1 0 1 1 0 1] ( 45 )
Input 2      : [0 0 0 1 1 0 1 0] ( 26 )
True         : [0 1 0 0 0 1 1 1] ( 71 )
Predicted    : [0 0 1 1 0 1 1 1] ( 55 )
Absolute error : 3.73035
X-entropy loss : 5.47834
-----
```

```
Input 1      : [0 0 0 0 1 1 1 0] ( 14 )
Input 2      : [0 0 1 1 0 1 0 1] ( 53 )
True         : [0 1 0 0 0 0 1 1] ( 67 )
Predicted    : [1 0 1 1 1 0 1 1] ( 187 )
Absolute error : 4.29031
X-entropy loss : 6.66568
-----
```

```

Input 1      : [0 0 1 0 0 0 0 1] ( 33 )
Input 2      : [0 1 0 1 1 1 1 0] ( 94 )
True         : [0 1 1 1 1 1 1 1] ( 127 )
Predicted    : [1 1 1 1 1 1 1 1] ( 255 )
Absolute error : 2.57187
X-entropy loss : 3.17627
-----

```

```

Input 1      : [0 0 1 1 0 0 1 0] ( 50 )
Input 2      : [0 0 0 0 0 0 1 1] ( 3 )
True         : [0 0 1 1 0 1 0 1] ( 53 )
Predicted    : [1 0 1 1 0 0 0 1] ( 177 )
Absolute error : 3.221
X-entropy loss : 4.36655
-----

```

```

Input 1      : [0 1 1 1 1 1 1 1] ( 127 )
Input 2      : [0 1 0 1 1 0 1 0] ( 90 )
True         : [1 1 0 1 1 0 0 1] ( 217 )
Predicted    : [1 0 1 0 0 1 0 1] ( 165 )
Absolute error : 4.30266
X-entropy loss : 6.77041
-----

```

```

Input 1      : [0 1 1 1 1 0 0 1] ( 121 )
Input 2      : [0 1 1 0 1 0 0 0] ( 104 )
True         : [1 1 1 0 0 0 0 1] ( 225 )
Predicted    : [1 0 0 1 0 0 0 1] ( 145 )
Absolute error : 3.78583
X-entropy loss : 5.52448
-----

```

```

Input 1      : [0 0 0 0 0 0 1 1] ( 3 )
Input 2      : [0 1 0 1 1 1 1 1] ( 95 )
True         : [0 1 1 0 0 0 1 0] ( 98 )
Predicted    : [1 0 0 1 1 1 0 0] ( 156 )
Absolute error : 4.98435
X-entropy loss : 8.13293
-----

```

```

Input 1      : [0 1 0 0 1 0 1 0] ( 74 )
Input 2      : [0 1 1 0 1 1 0 1] ( 109 )
True         : [1 0 1 1 0 1 1 1] ( 183 )
Predicted    : [1 0 1 0 0 1 1 1] ( 167 )
Absolute error : 2.79596
X-entropy loss : 3.56718
-----

```

```

Input 1      : [0 0 1 0 1 1 0 0] ( 44 )
Input 2      : [0 0 1 1 1 0 1 1] ( 59 )
True         : [0 1 1 0 0 1 1 1] ( 103 )
Predicted    : [1 0 0 1 0 1 1 1] ( 151 )
Absolute error : 3.66828
X-entropy loss : 5.39607
-----

```

```

Input 1      : [0 0 0 1 1 0 1 0] ( 26 )
Input 2      : [0 1 0 0 0 1 0 0] ( 68 )
True         : [0 1 0 1 1 1 1 0] ( 94 )
Predicted    : [1 0 0 1 1 1 1 0] ( 158 )
Absolute error : 3.09235
X-entropy loss : 4.10179
-----

```

```
In [ ]: sess.close()
```